

Framing Rules

Henry Blanchette

1 Definitions

Note: in this document “formula” refers to “precise formula,” however gradual formulas will eventually be supported.

A **permission** is to either access a field, written $\text{access}(e.f)$, or to assume a predicate holds of its arguments, written $\text{assume}(\alpha_C(\bar{e}))$.

A formula ϕ **requires** a permission π if ϕ contains an access or assumption that π permits. The set of all permissions that ϕ requires (the set of permissions required to frame ϕ) is called the **requirements** of ϕ .

A formula ϕ **grants** permission π if it contains an adjunct that yields π .

A set of permissions Π **frames** a formula ϕ if and only if ϕ requires only permissions contained in Π , written

$$\Pi \models_I \phi.$$

The **footprint** of a formula ϕ is the smallest permission mask that frames ϕ , written

$$[\phi].$$

A formula ϕ is **self-framing** if and only if for any set of permissions ϕ , $\Pi \models_I \phi$, written

$$\vdash_{\text{frm}I} \phi.$$

In other words, ϕ is self-framing if and only if it grants all the permissions that it requires.

2 Framing without Aliasing

For this section framing decisions do not consider aliasing, for the sake of an introduction.

2.1 Deciding Framing without Aliasing

The following algorithm decides $\Pi \models_I \phi$ for a given set of permissions Π and formula ϕ .

$\Pi \models_I \phi \iff$	match ϕ with	v, x	$\mapsto \top$
		$e_1 \oplus e_2$	$\mapsto \Pi \models_I e_1, e_2$
		$e_1 \odot e_2$	$\mapsto \Pi \models_I e_1, e_2$
		$e.f$	$\mapsto \Pi \models_I e \wedge \text{acc}(e.f) \in \Pi$
		$\text{acc}(e.f)$	$\mapsto \Pi \models_I e$
		$\phi_1 \circledast \phi_2$	$\mapsto \Pi \cup \text{granted}(\phi_1 \circledast \phi_2) \models_I \phi_1, \phi_2$
		$\alpha_C(\bar{e})$	$\mapsto \Pi \models_I \bar{e}$
		if e then ϕ_1 else ϕ_2	$\mapsto \Pi \models_I e, \phi_1, \phi_2$
		unfolding $\alpha_C(\bar{e})$ in ϕ	$\mapsto \text{assume}(\alpha_C(\bar{e})) \in \Pi \wedge \Pi \models_I \alpha_C(\bar{e}) \wedge \Pi \models_I \phi$

The following algorithm collects the set of permissions granted by a given formula ϕ .

$\text{granted}(\phi) :=$	match ϕ with	e	$\mapsto \emptyset$
		$\text{acc}(e.f)$	$\mapsto \{\text{access}(e.f)\}$
		$\phi_1 \circledast \phi_2$	$\mapsto \text{granted}(\phi_1) \cup \text{granted}(\phi_2)$
		$\alpha_C(\bar{e})$	$\mapsto \{\text{assume}(\alpha_C(\bar{e}))\}$
		if e then ϕ_1 else ϕ_2	$\mapsto \text{granted}(\phi_1) \cap \text{granted}(\phi_2)$
		unfolding $\alpha_C(\bar{e})$ in ϕ	$\mapsto \text{granted}(\phi)$

2.2 Notes

- The conditional expression e in a formula of the form (**if e then ϕ_1 else ϕ_2**) is considered indeteminant for the purposes of statically deciding framing.
- The body formula ϕ in a formula of the form (**unfolding $\text{acc}_C(\bar{e})$ in ϕ**) does not have to make use of the **assume**($\text{acc}_C(\bar{e})$) required by the structure.

2.3 Deciding Self-Framing without Aliasing

The following algorithm decides $\vdash_{\text{frm}I} \phi$ for a given formula ϕ .

$$\vdash_{\text{frm}I} \phi \iff \emptyset \models_I \phi$$

3 Aliasing (New)

Let I be the set of identifiers. A pair of identifiers $x, y \in I$ are **unique** if they are not the same identifier. The proposition that x, y are unique is written $\text{unique}(x, y)$ (note that this is importantly different notation from $x = y$). The proposition that two identifiers refer to the same memory in the heap is written $x = y$. A set of identifiers $\{x_\alpha\}$ is **aliasing** if and only if each x_α refers to the same memory in the heap i.e.

$$x_{\alpha_1} = \dots = x_{\alpha_k} \text{ where } \{\alpha\} = \{\alpha_1, \dots, \alpha_k\}.$$

The proposition that $\{x_\alpha\}$ is aliasing is written $\text{aliasing } \{x_\alpha\}$.

An **aliasing context** is a set A of aliasing propositions. As a set of propositions, the consistency of A can be considered. Explicitly, A is consistent if and only if

$$\nexists x, y \in I : \text{unique}(x, y) \wedge A \vdash \text{aliasing } \{x, y\} \wedge \sim \text{aliasing } \{x, y\}$$

The proposition that A is consistent is written $\text{consistent}(A)$.

An alias context is **overlapping** if and only if there exist at least two unique sets of identifiers such that they have a non-empty intersection and both are asserted aliasing in A i.e.

$$\exists I_1, I_2 \subset I : (I_1 \neq I_2) \wedge (I_1 \cap I_2 \neq \emptyset) \wedge (\text{aliasing}(I_1) \in A) \wedge (\text{aliasing}(I_2) \in A)$$

An overlapping alias context is inefficient for deciding the propositions that it entails. Fortunately the framing-deciding algorithm I present ensures that its tracked alias context never becomes overlapping.

A alias context A is **full** if and only if

$$\forall I_\alpha \subset I : A \vdash P(I_\alpha) \implies \exists P(I_{\alpha'}) \in A : I_\alpha \subset I_{\alpha'}$$

where P is an aliasing predicate (either aliasing or $\sim \text{aliasing}$). In other words, a full alias context is the most efficient representation of its total propositional strength. This is useful for efficient computation, as demonstrated in the following.

Given A an alias context and $x \in I$ an identifier, define:

$\text{aliases-of}(x) :=$ the largest set such that $x \in \text{aliases-of}(x) \wedge A \vdash \text{aliasing}(\text{aliases-of}(x))$

$\text{not-aliases-of}(x) :=$ the largest set such that $\forall x' \in \text{not-aliases-of}(x) : A \vdash \sim \text{aliasing } \{x, x'\}$

If A is non-overlapping and full, the computation of $\text{aliases-of}(x)$ is simply the extraction from A the proposition that asserts aliasing of a set of identifiers that contains x and the computation of $\text{not-aliases-of}(x)$ is the collection of all identifiers other than x mentioned in propositions of A that assert the negation of aliasing with x . For example,

$$A := \{\text{aliasing } \{x, y\}, \text{aliasing } \{z\}, \sim \text{aliasing } \{x, z\}, \sim \text{aliasing } \{y, z\}\}$$

id	$\text{aliases-of}(id)$	$\text{not-aliases-of}(id)$
x	$\{x, y\}$	$\{z\}$
y	$\{x, y\}$	$\{z\}$
z	$\{z\}$	$\{x, y\}$

4 Aliasing (Old)

The **alias status** of a set of identifiers $\{x_\alpha\}$ is exactly one of the following: **aliases**, **non-aliases**, **undetermined-aliases**.

- The x_α are **aliases** if each x_α refers to the same memory in the heap.
- The x_α are **non-aliases** if each x_α refers to distinct memory in the heap.
- The x_α are **undetermined-aliases** if they may be **aliases** or **non-aliases**.

An **alias class** is a pair $[S, I]$ where S is an alias status and I is a set of identifiers where the identifiers of I have alias status S . $\text{identifiers}(\{[S_\alpha, I_\alpha]\}) := \bigcup I_\alpha$ is the set of identifiers of a set of alias classes. It is possible to keep track of **undetermined-aliases** classes. However, for the sake of efficiency, some give identifiers are considered **undetermined-aliases** if no subset of them are asserted as **aliases** nor **non-aliases** by any alias class.

A set of alias classes $\{[S_\alpha, I_\alpha]\}$ is **overlapping** if and only if

$$\bigcup I_\alpha \neq \emptyset.$$

A set of alias statuses $\{S_\alpha\}$ is **compatible** if and only if

$$\forall S \in \{S_\alpha\} : \forall \alpha : S_\alpha = S$$

A set of alias classes A is **compatible** if and only if

$$\forall \{[S_\alpha, I_\alpha]\} \subset A : \{[S_\alpha, I_\alpha]\} \text{ is overlapping} \implies \{S_\alpha\} \text{ is compatible}$$

This is to say that a set of compatible alias classes must not assert that a pair of identifiers are both **aliases** and **non-aliases** — every overlapping set of alias classes is compatible.

Given two compatible sets of alias classes A, A' , the compatibility of $A \cup A'$ can be considered, written $A \uplus A'$. Deciding $A \uplus A'$ reduces to computing $\text{simplify}(A \cup A')$ which either preserves compatibility or raises an exception, where

$$\begin{aligned} \text{simplify}(\{[S_\alpha, I_\alpha]\}) &:= \left\{ \left[S, \bigcup I_{\alpha_i} \right] \mid \forall \alpha : \{\alpha_i\} = \text{LOS}(\alpha) \wedge ((\forall i : S = S_{\alpha_i}) \vee (\text{raise exception})) \right\}, \\ \text{LOS}(\alpha) &:= \{\alpha_i\}, \text{ the largest subset of } \{\alpha\} \\ &\quad \text{such that } \alpha \in \{\alpha_i\} \text{ and } \{I_{\alpha_i}\} \text{ is overlapping.} \end{aligned}$$

For each set of overlapping alias classes, simplify either combines them or throws an exception.

4.1 Deciding Alias Class Compatibility

A set of alias classes A is **compatible** with a formula ϕ if and only if A is compatible with the aliasing assertions yielded by ϕ , written $A \uplus \phi$. The following algorithm decides $A \uplus \phi$.

$A \uplus \phi \iff \text{match } \phi \text{ with}$	e	$\mapsto A \uplus \text{asserted}(e)$
	$\text{acc}(e.f)$	$\mapsto A \uplus \{[\text{non-aliases}, \{e\} \cup \text{identifiers}(A)]\}$
	$\phi_1 \otimes \phi_2$	$\mapsto A \uplus \text{asserted}(\phi_1) \uplus \text{asserted}(\phi_2)$
	$\phi_1 \wedge \phi_2$	$\mapsto (A \uplus \phi_1) \wedge (A \uplus \phi_2)$
	$\text{if } e \text{ then } \phi_1 \text{ else } \phi_2$	$\mapsto \text{if } A \uplus \text{asserted}(e)$ $\text{then } A \uplus \text{asserted}(e) \uplus \text{asserted}(\phi_1)$ $\text{else } A \uplus \neg \text{asserted}(e) \uplus \text{asserted}(\phi_2)$
	ϕ	$\mapsto A \uplus \text{asserted}(\phi)$

where

$$\begin{aligned} \neg A &:= \{[\neg S_\alpha, I_\alpha] \mid [S_\alpha, I_\alpha] \in A\}, \\ \neg \text{aliases} &:= \text{non-aliases}, \\ \neg \text{non-aliases} &:= \text{aliases}. \end{aligned}$$

The following algorithm collects the set of alias classes asserted by a given formula ϕ .

$\text{asserted}(\phi) := \text{match } \phi \text{ with}$	$x = y$	$\mapsto \{[\text{aliases}, \{x, y\}]\}$
	$x \odot y$	$\mapsto \{[\text{non-aliases}, \{e\}]\}$
	e	$\mapsto \emptyset$
	$\text{unfolding } \alpha(\bar{e}) \text{ in } \phi$	$\mapsto A \uplus \phi$

The following algorithm collects the set of identifiers in a given set of alias classes A .

$$\text{identifiers}(\{[S_\alpha, I_\alpha]\}) := \bigcup I_\alpha$$

5 Framing with Aliasing

For this section framing decisions *do* consider aliasing.

5.1 Deciding Framing with Aliasing

5.2 Notes

5.3 Deciding Self-Framing with Aliasing

The following algorithm decides $\vdash_{\text{frm}I} \phi$ for a given formula ϕ .